

# MLR

Brian Dehlinger

December 1, 2019

## Original Model Train Test Split, Boxplot

```
set.seed(10035)
remove_singularities <- function(dataset, gene_name){
  dataset_copy <- dataset
  item <- paste(gene_name, "~.", sep="")
  full_formula <- as.formula(item)
  fit <- lm(full_formula, data=dataset)
  singularities <- attributes(alias(fit)$Complete)$dimnames[[1]]
  for (singularity in singularities){
    dataset_copy[singularity] <- NULL
  }
  return(dataset_copy)
}
read_in_pruned_datasets_for_gene_0.8 <- function(gene_name, path){
  full_path0.8 <- paste(path, gene_name, "_for_r_0.8.txt", sep="")
  Data0.8 <- read.table(full_path0.8, header=TRUE, sep=',')
  Data0.8 <- remove_singularities(Data0.8, gene_name)
  return(Data0.8)
}
#install.packages("DAAG")
#install.packages("caret")
#install.packages("lmtest")
#install.packages("MASS")
#install.packages("car")
#install.packages("reshape")
#install.packages("plotmo")
#install.packages("olsrr")
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
library(lmtest)
```

```
## Loading required package: zoo
```

```
##
## Attaching package: 'zoo'
```

```
## The following objects are masked from 'package:base':
##
##   as.Date, as.Date.numeric
```

```
library(MASS)
library(car)
```

```
## Loading required package: carData
```

```
library(reshape)
library(plotmo)
```

```
## Loading required package: Formula
```

```
## Loading required package: plotrix
```

```
## Loading required package: TeachingDemos
```

```
library(olsrr)
```

```
##
## Attaching package: 'olsrr'
```

```
## The following object is masked from 'package:MASS':
##
##   cement
```

```
## The following object is masked from 'package:datasets':
##
##   rivers
```

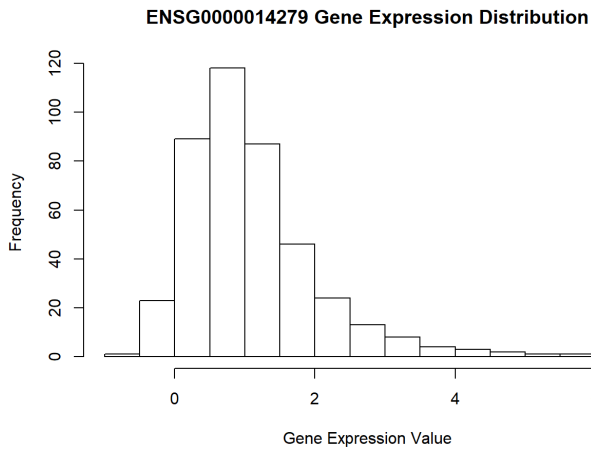
```
library(DAAG)
```

```
##
## Attaching package: 'DAAG'
```


```
## The following object is masked from 'package:car':
##
##   vif
```

```
## The following object is masked from 'package:MASS':
##
##   hills
```

```
gene_data <- read_in_pruned_datasets_for_gene_0.8("ENSG00000142794", "D:\\Project\\GitStash\\Applied_regression_pr
object\\")
gene_data <- as.data.frame(gene_data)
hist(gene_data[["ENSG00000142794"]], main="ENSG0000014279 Gene Expression Distribution", ylab="Frequency", xlab="
Gene Expression Value")
```



```
boxplot(gene_data[["ENSG00000142794"]], main="Boxplot of ENSG0000014279 expression", ylab="Expression value", xlab="Boxplot")
```



Boxplot

```
# We notice that the Expression Data is skewed to the right.
trainIndex <- createDataPartition(gene_data[["ENSG00000142794"]], p=.8, list = FALSE, times=1)

# We do test train split and explicitly create the data partition indexes
# because of the difference seed values produce in different R versions
```

```
gene_data <- gene_data[train_index,]
gene_test <- gene_data[~train_index,]
#We do a train validation split here.
```

## First Model Diagnostics

```
model <- lm(ENSG00000142794~., data=gene_data)
summary(model)
```


```
##
## Call:
## lm(formula = ENSG00000142794 ~ ., data = gene_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.4584 -0.3311 -0.0545  0.2781  3.2295
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1.0773438  0.8375846   1.286   0.1996
## rs12734589  -0.2426399  0.3650244  -0.665   0.5069
## rs2004380    0.1102502  1.0093512   0.109   0.9131
## rs10737456  -0.1014757  0.2853902  -0.356   0.7225
## rs10799692   0.0890538  0.5131390   0.174   0.8624
## rs3820296   -0.2600382  0.3517204  -0.739   0.4604
## rs10916988   0.1944314  0.1744793   1.114   0.2663
## rs1976403    0.5034704  0.4112980   1.224   0.2222
## rs1780324    0.4284627  0.2894724   1.480   0.1402
## rs4654753   -0.4658774  0.4593801  -1.014   0.3116
## rs79760554   0.0278784  0.1362679   0.205   0.8381
## rs12729540  -0.0386049  0.2499062  -0.154   0.8774
## rs12043777  -0.1137017  0.1445647  -0.787   0.4324
## rs10799702   0.1519789  0.2212533   0.687   0.4928
## rs4061838   -0.1683068  0.1846747  -0.911   0.3630
## rs7545635    0.0693701  0.2444405   0.284   0.7768
## rs4654946   -0.0593146  0.5275433  -0.112   0.9106
## rs2047653   -0.0215676  0.4805840  -0.045   0.9642
## rs4079441    0.0160076  0.5254259   0.030   0.9757
## rs12034222   0.1939562  0.1212573   1.600   0.1111
## rs3820292   -0.0718418  0.1268148  -0.567   0.5716
## rs75417796   0.1798876  0.3004644   0.599   0.5500
## rs7541779    0.0238033  0.2965717   0.080   0.9361
## rs75417790   0.0475355  0.2394576   0.199   0.8428
## rs10799691  -0.0251948  0.2950647  -0.085   0.9320
## rs1971328    0.2014015  0.2281684   0.883   0.3783
## rs12744514   0.6974904  0.5034685   1.385   0.1673
```

```

## rs41310392 -0.4924510 0.3108441 -1.584 0.1145
## rs75056920 0.0976996 0.1316812 0.742 0.4589
## rs1697421 -0.2674773 0.2314293 -1.156 0.2490
## rs146767219 0.0388611 0.1074221 0.362 0.7179
## rs12132412 0.0488435 0.1550308 0.315 0.7530
## rs112985962 -0.1747530 0.6483509 -0.270 0.7878
## rs147900768 0.1793372 0.2227159 0.805 0.4215
## rs1814737 -0.6478764 0.7304950 -0.887 0.3760
## rs56286426 -0.0108865 0.1338131 -0.081 0.9352
## rs17420195 1.2040546 0.9058701 1.329 0.1851
## rs113548640 0.1619209 0.3515875 0.461 0.6456
## rs146564096 -0.1459169 0.1311852 -1.112 0.2672
## rs10916989 0.0113086 0.2306912 0.049 0.9609
## rs150401191 0.0091174 0.6004537 0.015 0.9879
## rs3010180 -0.1346392 0.1276371 -1.055 0.2926
## rs78885464 0.1262377 0.1141166 1.106 0.2698
## rs6694671 0.0356021 0.3917338 0.091 0.9277
## rs144522615 -0.0696422 0.1296253 -0.537 0.5916
## rs61775934 -0.1628478 0.4170355 -0.390 0.6965
## rs3010181 0.1408603 0.1079256 1.305 0.1931
## rs1780323 -0.0641018 0.1257974 -0.510 0.6108
## rs61778366 -0.3494258 0.2960935 -1.180 0.2392
## rs79659018 0.0159360 0.1274369 0.125 0.9006
## rs10737458 -0.3630167 0.4858445 -0.747 0.4557
## rs11584744 0.5002325 0.3351833 1.492 0.1369
## rs201590042 -0.0149337 0.1038923 -0.144 0.8858
## rs1566524 -0.1873926 0.4552565 -0.412 0.6810
## rs61775950 0.0831508 0.0978564 0.850 0.3964
## rs144443608 0.0877285 0.1013082 0.866 0.3874
## rs10916986 0.3573069 0.2995042 1.193 0.2341
## rs904927 0.1363907 0.3373079 0.404 0.6863
## rs190849739 0.0054703 0.1267022 0.043 0.9656
## rs143561157 0.1553284 0.1161436 1.337 0.1824
## rs12740648 -0.0110182 0.1509227 -0.073 0.9419
## rs12118362 0.0390582 0.1495256 0.261 0.7942
## rs1809914 -0.4849652 0.2207155 -2.197 0.0290 *
## rs41310412 0.0082098 0.2822820 0.029 0.9768
## rs6423191 0.0007244 0.1983363 0.004 0.9971
## rs140609058 -0.2737547 0.3502827 -0.782 0.4353
## rs7555005 -0.1160461 0.2698248 -0.430 0.6675
## rs71512991 -0.0922331 0.1062848 -0.868 0.3864
## rs139043162 -0.0988816 0.3860193 -0.256 0.7981
## rs2800935 -0.0890408 0.1292364 -0.689 0.4915
## rs113324018 0.1567611 0.2392068 0.655 0.5129
## rs904928 0.0440956 0.1553133 0.284 0.7767
## rs7533048 0.1191272 0.2959662 0.403 0.6877
## rs12239666 -0.1800842 0.2227685 -0.808 0.4197
## rs972662 0.1707553 0.2390314 0.714 0.4757
## rs9726624 -0.0631679 0.1772911 -0.356 0.7219
## rs2800774 0.0266631 0.0939158 0.284 0.7767
## rs56268398 0.0605925 0.1983985 0.305 0.7603
## rs12082914 0.1894957 0.3400845 0.557 0.5779
## rs78566499 0.0775510 0.2753531 0.282 0.7785
## rs1809915 0.2982296 0.1990040 1.499 0.1353
## rs7547671 0.2785216 0.8860444 0.314 0.7535
## rs146549873 -0.0763266 0.0938884 -0.813 0.4171
## rs113934925 -0.1689566 0.9265713 -0.182 0.8555
## rs200384063 -0.0353196 0.2725797 -0.130 0.8970
## rs4654930 0.3933332 0.3124555 1.259 0.2093
## rs3855556 0.0959987 0.1545367 0.621 0.5351
## rs35836191 -0.6655432 0.5295844 -1.257 0.2101
## rs60803995 0.1947940 0.1536928 1.267 0.2063
## rs12037596 -0.1170011 0.1806859 -0.648 0.5179
## rs2651402 0.4172604 0.3110401 1.342 0.1811
## rs41290414 -0.3163552 0.4893582 -0.646 0.5186
## rs72476502 0.0790231 0.2649570 0.298 0.7658
## rs58090121 -0.0165627 0.1430192 -0.116 0.9079
## rs904937 -0.1079859 0.1691946 -0.638 0.5239
## rs41265985 -0.3857862 0.2618278 -1.473 0.1420
## rs77025042 -0.0157810 0.1191521 -0.132 0.8947
## rs7417849 -0.1141053 0.1663477 -0.686 0.4934
## rs113558389 -0.1506948 0.2873036 -0.525 0.6004
## rs904929 0.1129374 0.1973569 0.572 0.5677
## rs10799704 -0.0001877 0.0840813 -0.002 0.9982
## rs148061397 0.0406450 0.2051900 0.198 0.8432
## rs115933091 0.1055186 0.3592554 0.294 0.7692
## rs143553186 0.5298658 0.2947059 1.798 0.0735 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6711 on 233 degrees of freedom
## Multiple R-squared:  0.6148, Adjusted R-squared:  0.4446
## F-statistic: 3.611 on 103 and 233 DF,  p-value: 2.972e-16

```

```
plot(model, which=1)
```




```
bptest(model)
```

```
##
## studentized Breusch-Pagan test
##
## data: model
## BP = 96.199, df = 103, p-value = 0.6693
```

```
# We note from the Residuals Vs Fitted Plot that there
# might be a slight concern with nonconstant variance but for the Breusch-Pagan test
# we fail to reject with an alpha of 0.05.
```

```
plot(model, which=2)
```




```
shapiro.test(model$residuals)
```

```
##
## Shapiro-Wilk normality test
##
## data: model$residuals
## W = 0.94027, p-value = 2.08e-10
```


```
# However, we notice an issue with normality from a QQPlot that
# suggests the data is skewed and we reject the shapiro-wilks test with an alpha of 0.05
```

```
plot(model, which=3)
```




```
#The Scale-Location plot has a slope most likely because there isn't enough data for the fitted values.
```

```
plot(model, which=4)
```



```
plot(model, which=5)
```




```
# We don't notice an issues with overly influential points in the residuals vs Leverage Plot.  
# However, we do note some points are considered to have high leverage later on. But these points  
# are actually important to the variation we want to capture in our data.
```

```
summary(model)$adj.r.squared
```

```
## [1] 0.444552
```

```
# We have an adjusted R-squared of 0.44452.
```

```
test <- model <- lm(ENSG00000142794+2~., data=gene_data)  
# We do a Boxcox on the model and note that we should do a transformation on Y(1/sqrt(Y+2))  
boxcox(test)
```



VIFs of Original Model

```
# VIF Calculation and setting N and P
n <- nrow(gene_data)
p <- length(model$coefficients)

# There is an indication of severe multicollinearity.
car::vif(model)
```

```
## rs12734589 rs2004380 rs10737456 rs10799692 rs3820296 rs10916988
## 53.938076 396.667047 29.989972 98.776324 41.200683 11.627147
## rs1976403 rs1780324 rs4654753 rs79760554 rs12729540 rs12043777
## 59.105285 32.812037 70.967745 5.029291 22.203345 8.439732
## rs10799702 rs4061838 rs7545635 rs4654946 rs2047653 rs4079441
## 18.192176 12.458436 21.141962 85.906720 73.463714 89.795592
## rs12034222 rs3820292 rs75417796 rs7541779 rs75417790 rs10799691
## 6.065616 6.054684 1.756140 27.444465 1.846154 27.469175
## rs1971328 rs12744514 rs41310392 rs75056920 rs1697421 rs146767219
## 15.900064 103.196154 40.269403 7.224826 23.143299 3.799889
## rs12132412 rs112985962 rs147900768 rs1814737 rs56286426 rs17420195
## 8.346162 113.962670 8.909498 131.440962 7.271876 199.760324
## rs113548640 rs146564096 rs10916989 rs150401191 rs3010180 rs78885464
## 30.700811 5.847086 19.261110 94.030738 8.797080 4.748219
## rs6694671 rs144522615 rs61775934 rs3010181 rs1780323 rs61778366
## 46.954267 8.112206 53.300421 6.207780 4.287989 20.657987
## rs79659018 rs10737458 rs11584744 rs201590042 rs1566524 rs61775950
## 8.109858 65.126021 31.466721 3.403908 77.820080 3.513458
## rs144443608 rs10916986 rs904927 rs190849739 rs143561157 rs12740648
## 2.019597 33.746042 28.487079 9.052471 6.058666 4.889190
## rs12118362 rs1809914 rs41310412 rs6423191 rs140609058 rs7555005
## 5.767630 14.932888 7.027638 17.147129 17.061526 27.146524
## rs71512991 rs139043162 rs2800935 rs113324018 rs904928 rs7533048
## 2.913536 16.995030 8.307933 11.325202 4.666285 17.181528
## rs12239666 rs972662 rs9726624 rs2800774 rs56268398 rs12082914
## 14.776399 13.066115 2.152436 2.108870 7.648545 30.252107
## rs78566499 rs1809915 rs7547671 rs146549873 rs113934925 rs200384063
## 9.431760 13.113966 240.217181 3.809010 144.059295 4.097976
## rs4654930 rs3855556 rs35836191 rs60803995 rs12037596 rs2651402
## 19.221395 8.404156 35.165796 8.655057 12.487565 1.881941
## rs41290414 rs72476502 rs58090121 rs904937 rs41265985 rs77025042
## 37.250580 9.801619 3.754337 9.915960 12.999266 1.703049
## rs7417849 rs113558389 rs904929 rs10799704 rs148061397 rs115933091
## 8.494297 8.513513 12.096988 1.916642 1.841093 1.689051
## rs143553186
## 1.506332
```

```
test <- lm(ENSG00000142794+2~., data=gene_data)
```

## Doing the Added Variable Plots


```
# The added variable plots suggest a lot of the predictors do not
# add new information when the other predictors are included in the model.
pdf("AddedVariablesBeforeSelection.pdf")
avPlots(model, ask=FALSE)
dev.off()
```

```
## png
## 2
```

## Transforming the Data

```
# We do the transform on the train and test datasets. # We also set n.
# Additionally, we also plot the transformed values and note they appear to be normally distributed.
transformed_gene_data <- gene_data
transformed_gene_test <- gene_test
n <- nrow(transformed_gene_data)
transformed_gene_data[["ENSG00000142794"]] <- 1/(sqrt(transformed_gene_data[["ENSG00000142794"]]+2))
transformed_gene_test[["ENSG00000142794"]] <- 1/(sqrt(transformed_gene_test[["ENSG00000142794"]]+2))
boxplot(transformed_gene_data[["ENSG00000142794"]], main="Boxplot of Transformed Gene Expression Values", ylab="Gene
Expression Transformed Value", xlab="Boxplot")
```

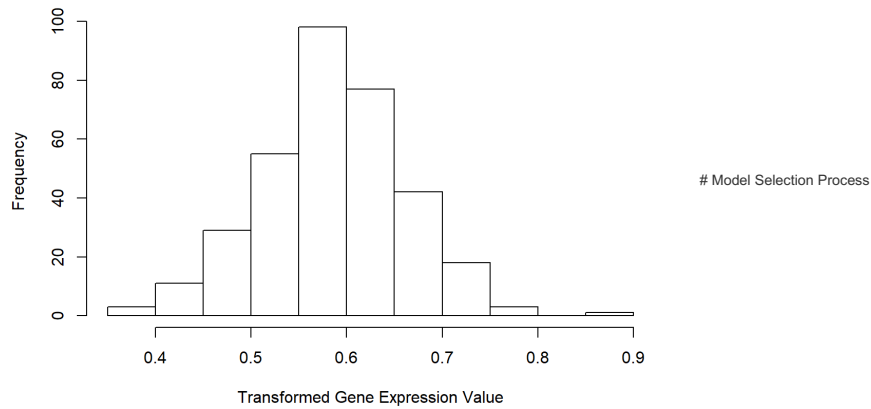
Boxplot of Transformed Gene Expression Values



Boxplot

```
hist(transformed_gene_data[["ENSG00000142794"]], main="Transformed ENSG0000014279 Gene Expression Distribution", y
lab="Frequency", xlab="Transformed Gene Expression Value")
```

## Transformed ENSG0000014279 Gene Expression Distribution



```
# We do variable selection and we do a strict variable selection with a p-value of 0.01. Step-Forward BIC, and Step-Backward BIC.
# We then put the selected variables for each model together and do a best subset exhaustive search over just these variables.

naive_model <- lm(ENSG0000014279~., data=transformed_gene_data)
naive_model_empty <- lm(ENSG0000014279~1, data=transformed_gene_data)

selected_p <- ols_step_backward_p(naive_model, prem=0.01)
selected_p_model <- selected_p$model
p <- length(selected_p_model$coefficients)
RSS_selected <- c(crossprod(selected_p_model$residuals))
MSE <- RSS_selected / length(selected_p_model$residuals)
selected_p_stats <- rbind(sqrt(MSE), BIC(selected_p_model), summary(selected_p_model)$adj.r.squared, DAAG::press(selected_p_model))

selected<- stepAIC(naive_model, k=log(n))
p <- length(selected$coefficients)
RSS_selected <- c(crossprod(selected$residuals))
MSE <- RSS_selected / length(selected$residuals)
selected_stats <- rbind(sqrt(MSE), BIC(selected), summary(selected)$adj.r.squared, DAAG::press(selected))

forward_bic <- stepAIC(naive_model_empty, scope = list(upper=naive_model, lower=naive_model_empty), direction="forward", k=log(n))
p <- length(forward_bic$coefficients)
RSS_selected <- c(crossprod(forward_bic$residuals))
MSE <- RSS_selected / length(forward_bic$residuals)
forward_bic_stats <- rbind(sqrt(MSE), BIC(forward_bic), summary(forward_bic)$adj.r.squared, DAAG::press(forward_bic))
```

## VIF Stats of Selected Models

```
car::vif(selected)
```

```
## rs12734589 rs10799692 rs113548640 rs6694671 rs1566524 rs35836191
## 5.036327 4.997166 2.697843 3.118706 3.069425 1.694852
## rs60803995
## 2.351035
```

```
car::vif(selected_p_model)
```

```
## rs12734589 rs113548640 rs6694671 rs1566524 rs35836191 rs60803995
## 3.283511 2.146153 2.669397 3.068997 1.694835 2.347506
```

```
car::vif(forward_bic)
```

```
## rs12734589 rs1976403 rs4654753 rs6694671 rs113548640 rs10916989
## 3.192661 2.848736 4.706859 3.299466 5.526336 3.981726
```

```
# We now do not see any issues with multicollinearity which suggests
# we do not need to use ridge regression on these reduced variable models.
```

## Selected Model Choosing

```
search_full_model <- lm(ENSG0000014279~rs12734589 + rs1976403 + rs4654753 + rs6694671 + rs113548640 + rs10916989
+ rs1566524, rs35836191 + rs60803995 + rs10799692, data=transformed_gene_data)

all_possible <- ols_step_best_subset(search_full_model)
all_possible
```

```
## Best Subsets Regression
## -----
## Model Index Predictors
## -----
## 1 rs12734589
## 2 rs6694671 rs113548640
## 3 rs12734589 rs6694671 rs113548640
## 4 rs12734589 rs1976403 rs6694671 rs113548640
## 5 rs12734589 rs4654753 rs6694671 rs113548640 rs10916989
## 6 rs12734589 rs1976403 rs4654753 rs6694671 rs113548640 rs10916989
## 7 rs12734589 rs1976403 rs4654753 rs6694671 rs113548640 rs10916989 rs1566524
## -----
## Subsets Regression Summary
## -----
##
## Adj. Pred
## Model R-Square R-Square R-Square C(p) AIC SBIC SBC MSEP FPE
HSP APC
## -----
## 1 0.4428 0.4412 0.4356 12886.6397 -972.0540 NA -960.5938 0.0033 0.00
33 0.0000 0.5638
## 2 0.5173 0.5144 0.5077 11122.3098 -1018.3882 NA -1003.1079 0.0028 0.00
28 0.0000 0.4914
## 3 0.5479 0.5438 0.5357 10397.7557 -1038.4744 NA -1019.3740 0.0027 0.00
27 0.0000 0.4630
## 4 0.5550 0.5496 0.5398 10231.6838 -1041.7965 NA -1018.8760 0.0026 0.00
26 0.0000 0.4584
## 5 0.5639 0.5573 0.5477 10021.9716 -1046.6223 NA -1019.8817 0.0026 0.00
26 0.0000 0.4519
## 6 0.5751 0.5674 0.556 9758.9045 -1053.3680 NA -1022.8074 0.0026 0.00
26 0.0000 0.4429
## 7 0.5798 0.5709 0.5588 9649.2532 -1055.1209 NA -1020.7402 0.0025 0.00
25 0.0000 0.4406
## -----
## AIC: Akaike Information Criteria
## SBIC: Sawa's Bayesian Information Criteria
## SBC: Schwarz Bayesian Criteria
## MSEP: Estimated error of prediction, assuming multivariate normality
## FPE: Final Prediction Error
## HSP: Hocking's Sp
## APC: Amemiya Prediction Criteria

final_stats <- cbind(selected_p_stats,selected_stats,forward_bic_stats)
row.names(final_stats) <- c("RMSE", "BIC", "ADJ-R-Squared", "PRESS")
colnames(final_stats) <- c("BackwardSelected_By_P_0.01", "Backward_BIC", "ForwardBIC")
knitr::kable(final_stats)
```

	BackwardSelected_By_P_0.01	Backward_BIC	ForwardBIC
RMSE	0.0498458	0.0493727	0.0495121
BIC	-1018.2806629	-1018.8876841	-1022.8073740
ADJ-R-Squared	0.5615122	0.5684880	0.5673627
PRESS	0.8735692	0.8599983	0.8632144

```
# We note that the optimal model is the one selected by forward_bic
# if we consider all the variables all these models give us and their best possible subsets.
# We try to minimize SBC(BIC) and keep a reasonable Adj-R-Squared(>0.55) and MSE.
# We note that MSE is near minmal at 4 predictors. But we have a lower SBC with 6 predictors and a slightly higher R-squared.
model <- forward_bic
```

## Diagnostics of Chosen Model


```
# We redo the model diagnostics from the beginning on the selected model.
# We will then compare this to one of the other models in our testing.
summary(model)
```

```
##
## Call:
## lm(formula = ENSG00000142794 ~ rs12734589 + rs1976403 + rs4654753 +
## rs6694671 + rs113548640 + rs10916989, data = transformed_gene_data)
##
## Residuals:
## Min 1Q Median 3Q Max
## -0.134598 -0.026933 0.000069 0.028130 0.194233
##
## Coefficients:
## Estimate Std. Error t value Pr(>|t|)
## (Intercept) 0.524253 0.010447 50.183 < 2e-16 ***
## rs12734589 0.027072 0.006622 4.088 5.46e-05 ***
## rs1976403 -0.019831 0.006733 -2.946 0.003453 **
## rs4654753 0.033980 0.008821 3.852 0.000141 ***
## rs6694671 0.035719 0.007743 4.613 5.68e-06 ***
## rs113548640 -0.036296 0.011122 -3.263 0.001216 **
## rs10916989 -0.024367 0.007821 -3.116 0.001995 **
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.05003 on 330 degrees of freedom
## Multiple R-squared: 0.5751, Adjusted R-squared: 0.5674
## F-statistic: 74.44 on 6 and 330 DF, p-value: < 2.2e-16
```


```
# All variables appear significant

# We note that the residual vs fitted plot appears to be random.
plot(model, which=1)
```





```
# We notice a small issue with normality still but this can probably be ignored.
# There is no issue with non constant error variance which is more important.
plot(model, which=2)
```




```
shapiro.test(model$residuals)
```

```
##
## Shapiro-Wilk normality test
##
## data: model$residuals
## W = 0.98986, p-value = 0.01963
```


```
bptest(model)
```

```
##
## studentized Breusch-Pagan test
##
## data: model
## BP = 2.7847, df = 6, p-value = 0.8353
```


```
# We notice potential outliers in the Cook's Distance Plot and Residuals vs Leverage Plot
# so we will do further analysis on these points.
# Particular points 283, 309, and 317 have a high cook's distance and are candidate outliers.
# The Scale Location Plot doesn't appear to have a slope which is good.
plot(model, which=3)
```



```
plot(model, which=4)
```



```
plot(model, which=5)
```



```
# The added variable plots look much better.
```

```
pdf("AddedVariablesAfterSelection.pdf")
avPlots(model, ask=FALSE)
dev.off()
```

```
## png
## 2
```

```
# All variables appear significant.
# MSE looks very good.
knitr::kable(anova(model))
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
rs12734589	1	0.8610066	0.8610066	343.928080	0.0000000
rs1976403	1	0.1251881	0.1251881	50.006228	0.0000000
rs4654753	1	0.0464546	0.0464546	18.556256	0.0000218
rs6694671	1	0.0403177	0.0403177	16.104843	0.0000742
rs113548640	1	0.0208500	0.0208500	8.328506	0.0041598
rs10916989	1	0.0243041	0.0243041	9.708235	0.0019954
Residuals	330	0.8261383	0.0025034	NA	NA
# Sensitivity	Analys	is			

```
# We do outlier analysis here

cd <- cooks.distance(model)
transformed_gene_dropped <- model$model[-c(which(cd > 4 / (n-p))),]
cooks_distance_model <- lm(ENSG00000142794~., data=transformed_gene_dropped)

# We note that we would like to retain these points regardless as they are interesting and represent variation
# in the population that can be informative in clinical situations.
# We opt to not use Robust Regression since the effect on the final model is not serious enough.

leverage_threshold <- (3*(p))/n
leverage_values <- hatvalues(model)
transformed_gene_dropped_leverage <- model$model[-c(which(leverage_values > leverage_threshold)),]
leverage_model <- lm(ENSG00000142794~., data=transformed_gene_dropped_leverage)

betas <- as.data.frame(dfbetas(model))

# We print the largest DFBeta for each parameter and the index of that DFBeta.
to_remove <- c()
for(column in betas){
  outliers <- which(abs(column) > 0.11)
  print(max(abs(column)))
  print(which(abs(column) == max(abs(column))))
  to_remove <- union(to_remove, outliers)
}

## [1] 0.4216956
## [1] 238
## [1] 0.5156542
## [1] 245
## [1] 0.8014485
## [1] 221
## [1] 0.3773799
## [1] 221
## [1] 0.7861822
## [1] 221
## [1] 0.8713133
## [1] 221
## [1] 0.3917799
## [1] 238

# One of 238, 245, and 221 give the max DFBetas the coefficients. These are the greatest outliers for DFBetas.
transformed_gene_dropped_betas <- model$model[-c(to_remove),]
new_model_betas <- lm(ENSG00000142794~., data=transformed_gene_dropped_betas)
coefficients <- rbind(model$coefficients, new_model_betas$coefficients, cooks_distance_model$coefficients, leverage_model$coefficients)
row.names(coefficients) <- c("Normal Model", "DroppedHighDFBetas", "DroppedHighCooks", "DroppedHighLeverage")
```

## Sensitvity Analysis Continued

```
# We make a summary table that shows what happens when we drop high DFBetas, highCooks, and High Leverage Points from the model.
model_data <- c()
add_to_coefficients_data <- function(model, model_data){
  current_model <- model
  RSS_selected <- c(crossprod(current_model$residuals))
  MSE <- RSS_selected / length(current_model$residuals)
  current_model_stats <- cbind(sqrt(MSE), BIC(current_model), summary(current_model)$adj.r.squared, DAAG::press(current_model), shapiro.test(model$residuals)$p.value, bptest(model)$p.value)
  colnames(current_model_stats) <- c("RMSE", "BIC", "ADJ-R-Squared", "PRESS", "Shapiro_Wilks_Test_P_value", "Brecht-Pagan_Test_P_value")
  model_data <- rbind(model_data, current_model_stats)
  shapiro.test(model$residuals)
  return(model_data)
}
model_data <- add_to_coefficients_data(model, model_data)
model_data <- add_to_coefficients_data(cooks_distance_model, model_data)
model_data <- add_to_coefficients_data(new_model_betas, model_data)
model_data <- add_to_coefficients_data(leverage_model, model_data)

# Here is the summary table
summary_of_excluding_points <- cbind(coefficients, model_data)
knitr::kable(summary_of_excluding_points)
```

	(Intercept)	rs12734589	rs1976403	rs4654753	rs6694671	rs113548640	rs10916989	RMSE	BIC	ADJ-R-Squared	PRESS	Shapiro_Wilks_Te
Normal Model	0.5242534	0.0270721	-0.0198311	-0.0339801	0.0357185	-0.0362962	-0.0243674	0.0495121	1022.8074	-0.5673627	0.8632144	
DroppedHighDFBetas	0.5242154	0.0268792	-0.0159879	-0.0372453	0.0372679	-0.0397645	-0.0286545	0.0441744	1038.9073	-0.6071241	0.6451452	
DroppedHighCooks	0.5288562	0.0291684	-0.0170019	-0.0298446	0.0345575	-0.0387295	-0.0245814	0.0434263	-947.5295	0.6313659	0.5618019	
DroppedHighLeverage	0.5282415	0.0369003	-0.0150414	-0.0312738	0.0309058	-0.0392810	-0.0287681	0.0488392	-981.2104	0.5852552	0.7975638	

```
# We only would consider dropping HighDF Betas since dropping the other elements hurts BIC and might
# even cause the non constant error varirance assumption not to hold. We choose not to drop any points.
#The coefficients are not changed greatly enough to warrant us to exclude so many
# points that contain valuable variation from the overall population.

# Dropping the High DFBetas points might be something we would consider if
# we had more data to represent the true diversity of the population but dropping
# these points might exclude points that include important variation.
```

## Training the Final Model and Diagnostics of Final Model

```
# We compare the validation model and original model and get MSPR and compare it to MSE.

predictions <- predict(model, transformed_gene_test)
validation_model_one <- lm(formula(model), data=transformed_gene_test)

# Comparison of Model Stats
model_stats <- cbind(sigma(model)^2, summary(model)$adj.r.squared, DAAG::press(model))
validation_model_stats <- cbind(sigma(validation_model_one)^2, summary(validation_model_one)$adj.r.squared, DAAG::press(validation_model_one))
comparison_of_stats <- rbind(model_stats, validation_model_stats)
colnames(comparison_of_stats) <- c("MSE", "Adj-R-Squared", "PRESS")
rownames(comparison_of_stats) <- c("Original Model", "Validation Model")
comparison_of_stats
```

```
##              MSE Adj-R-Squared    PRESS
## Original Model   0.002503449    0.5673627 0.8632144
## Validation Model 0.002299010    0.5480023 0.1912980
```

```
knitr::kable(comparison_of_stats)
```

	MSE	Adj-R-Squared	PRESS
Original Model	0.0025034	0.5673627	0.8632144
Validation Model	0.0022990	0.5480023	0.1912980

```
# MSE vs MSPR
mspr_vs_mse <- cbind(sigma(model)^2, mean((transformed_gene_test[["ENSG00000142794"]]-predictions)^2))
colnames(mspr_vs_mse) <- c("MSE", "MSPR")
knitr::kable(mspr_vs_mse)
```

	MSE	MSPR
	0.0025034	0.0021703


```
# Fitting the Final Multiple Linear Regression Model :)
final_gene_data <- read_in_pruned_datasets_for_gene_0.8("ENSG00000142794", "D:\\Project\\GitStash\\Applied_regression_project\\")
transformed_final_gene_data <- final_gene_data
transformed_final_gene_data[["ENSG00000142794"]] <- 1/(sqrt(transformed_final_gene_data[["ENSG00000142794"]]+2))
final_model <- lm(formula(model), data=transformed_gene_data)
summary(final_model)
```

```
##
## Call:
## lm(formula = formula(model), data = transformed_gene_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.134598 -0.026933  0.000069  0.028130  0.194233
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.524253   0.010447  50.183  < 2e-16 ***
## rs12734589   0.027072   0.006622   4.088 5.46e-05 ***
## rs1976403   -0.019831   0.006733  -2.946 0.003453 **
## rs4654753    0.033980   0.008821   3.852 0.000141 ***
## rs6694671    0.035719   0.007743   4.613 5.68e-06 ***
## rs113548640 -0.036296   0.011122  -3.263 0.001216 **
## rs10916989  -0.024367   0.007821  -3.116 0.001995 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.05003 on 330 degrees of freedom
## Multiple R-squared:  0.5751, Adjusted R-squared:  0.5674
## F-statistic: 74.44 on 6 and 330 DF,  p-value: < 2.2e-16
```


```
knitr::kable(anova(final_model))
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
rs12734589	1	0.8610066	0.8610066	343.928080	0.0000000
rs1976403	1	0.1251881	0.1251881	50.006228	0.0000000
rs4654753	1	0.0464546	0.0464546	18.556256	0.0000218
rs6694671	1	0.0403177	0.0403177	16.104843	0.0000742
rs113548640	1	0.0208500	0.0208500	8.328506	0.0041598
rs10916989	1	0.0243041	0.0243041	9.708235	0.0019954
Residuals	330	0.8261383	0.0025034	NA	NA


```
plot(final_model, which=1)
```




```
plot(final_model, which=2)
```




```
plot(final_model, which=3)
```



```
plot(final_model, which=4)
```



```
plot(final_model, which=5)
```



```
shapiro.test(model$residuals)
```

```
##
## Shapiro-Wilk normality test
##
## data:  model$residuals
## W = 0.98986, p-value = 0.01963
```

```
bptest(final_model)
```

```
##
## studentized Breusch-Pagan test
##
## data:  final_model
## BP = 2.7847, df = 6, p-value = 0.8353
```

```
# We note the same diagnostic findings as we did from the selected model. There is a slight issue with normality and some outliers.
```

## Outlier Analysis in Final Model

```
n <- nrow(transformed_final_gene_data)
p <- length(final_model$coefficients)
model <- final_model
# We do outlier analysis here

cd <- cooks.distance(model)
transformed_gene_dropped <- model$model[-c(which(cd > 4 / (n-p))),]
cooks_distance_model <- lm(ENSG00000142794~., data=transformed_gene_dropped)

leverage_threshold <- (3*p)/n
leverage_values <- hatvalues(model)
transformed_gene_dropped_leverage <- model$model[-c(which(leverage_values > leverage_threshold)),]
leverage_model <- lm(ENSG00000142794~., data=transformed_gene_dropped_leverage)

betas <- as.data.frame(dfbetas(model))

# We print the largest DFBeta for each parameter and the index of that DFBeta.
to_remove <- c()
for(column in betas){
  outliers <- which(abs(column) > 2/sqrt(n))
  print(max(abs(column)))
  print(which(abs(column) == max(abs(column))))
  to_remove <- union(to_remove, outliers)
}
```

```
## [1] 0.4216956
## [1] 238
## [1] 0.5156542
## [1] 245
## [1] 0.8014485
## [1] 221
## [1] 0.3773799
## [1] 221
## [1] 0.7861822
## [1] 221
## [1] 0.8713133
## [1] 221
## [1] 0.3917799
## [1] 238
```

```
#238, 221, and 245 are possible massive outliers
transformed_gene_dropped_betas <- model$model[-c(to_remove),]
new_model_betas <- lm(ENSG00000142794~., data=transformed_gene_dropped_betas)
coefficients <- rbind(model$coefficients,new_model_betas$coefficients, cooks_distance_model$coefficients, leverage_model$coefficients)
row.names(coefficients) <- c("Normal Model", "DroppedHighDFBetas", "DroppedHighCooks", "DroppedHighLeverage")
```

## Outlier Analysis Continued of Final Model

```
# We make a summary table that shows what happens when we drop high DFBetas, highCooks, and High Leverage Points from the model.
model_data <- c()
add_to_coefficients_data <- function(model, model_data){
  current_model <- model
  RSS_selected <- c(crossprod(current_model$residuals))
  MSE <- RSS_selected / length(current_model$residuals)
  current_model_stats <- cbind(sqrt(MSE), BIC(current_model), summary(current_model)$adj.r.squared, DAAG::press(current_model), shapiro.test(model$residuals)$p.value, bptest(model)$p.value)
  colnames(current_model_stats) <- c("RMSE", "BIC", "ADJ-R-Squared", "PRESS", "Shapiro_Wilks_Test_P_value", "Brecht-Pagan_Test_P_value")
  model_data <- rbind(model_data, current_model_stats)
  shapiro.test(model$residuals)
  return(model_data)
}
model_data <- add_to_coefficients_data(model, model_data)
model_data <- add_to_coefficients_data(cooks_distance_model, model_data)
model_data <- add_to_coefficients_data(new_model_betas, model_data)
model_data <- add_to_coefficients_data(leverage_model, model_data)


# Here is the summary table
summary_of_excluding_points <- cbind(coefficients, model_data)
knitr::kable(summary_of_excluding_points)
```

	(Intercept)	rs12734589	rs1976403	rs4654753	rs6694671	rs113548640	rs10916989	RMSE	BIC	ADJ-R-Squared	PRESS	Shapiro_Wilks_Test_P_value
Normal Model	0.5242534	0.0270721	-0.0339801	0.0357185	-0.0362962	-0.0243674	0.0495121	1022.8074	-0.5673627	0.8632144		
DroppedHighDFBetas	0.5244619	0.0294151	-0.0376958	0.0333010	-0.0355336	-0.0282190	0.0434386	1029.1651	-0.6155703	0.6111564		
DroppedHighCooks	0.5242282	0.0298767	-0.0281422	0.0377016	-0.0447463	-0.0243940	0.0425174	-921.8093	0.6437677	0.5179535		
DroppedHighLeverage	0.5282157	0.0376963	-0.0327219	0.0280078	-0.0343316	-0.0281172	0.0490726	-952.8234	0.5846780	0.7864094		

```
#We note that dropping the high DFBetas again results in a satisfied normality assumption but we choose to retain these points.
```

## Leave One Out Cross Validation MSE of Final Model

```
p <- length(model$coefficients)
n <- nrow(model$model)
cv.lm(model, data=transformed_final_gene_data, m=n, printit = FALSE)
```



```
leave_one_out_mse <- press(model)/(n-p)
# Leave one out cross-validated MSE is 0.00262. This is excellent
```